

# MultiCloud, DevOps & AI Challenge - Day 2

## Part 1 – Docker

### Step 1: Install Docker on EC2

- Execute the following commands:

```
sudo yum update -y
sudo yum install docker -y
sudo systemctl start docker
sudo docker run hello-world
sudo systemctl enable docker
docker --version
sudo usermod -a -G docker $(whoami)
newgrp docker
```

- **sudo yum install docker -y** -> sudo → Runs the command as the superuser (administrator).
- yum → The package manager used in RHEL-based Linux distributions.
- update → Updates all system packages to their latest versions.
- -y → Automatically confirms "yes" to all prompts.
- **sudo yum install docker** - Installs Docker from the official package repository and the -y flag confirms installation without prompting the user.
- **sudo systemctl start docker** - systemctl → A command to control system services.
- start docker → Starts the Docker service.
- **sudo docker run hello-world** -> docker run hello-world - Runs a simple test container to verify that Docker is working.
- This pulls a small test image (hello-world) from Docker Hub and runs it.
- If Docker is working, it will print a welcome message.
- **sudo systemctl enable docker** - Ensures Docker starts automatically when the system boots up.
- **docker --version** - Displays the installed Docker version.
- **sudo usermod -a -G docker \$(whoami)** -> usermod -a -G docker \$(whoami) - usermod - Modify a user account.
- -a - Append the user to a group.
- -G docker - Add the user to the docker group.
- \$(whoami) - Gets the currently logged-in user.
- This allows the user to run Docker commands without sudo.

- **newgrp docker** -> Activates the new group membership without logging out.
- If you skip this step, you might need to log out and log back in for the changes to take effect.

## Step 2: Create Docker image for CloudMart

### Backend

- Create folder and download source code:
  - `mkdir -p challenge-day2/backend && cd challenge-day2/backend` - Creates a folder & Navigates into the backend directory.
  - `wget https://tcb-public-events.s3.amazonaws.com/mdac/resources/day2/cloudmart-backend.zip` - Command-line tool to download files from the internet.
  - `unzip cloudmart-backend.zip` - Extracts the contents of cloudmart-backend.zip into the current directory (backend/).
- Create .env file:
  - `vi .env`
  - Contents of .env:

```
PORT=5000
AWS_REGION=us-east-1
BEDROCK_AGENT_ID=<your-bedrock-agent-id> BEDROCK_AGENT_ALIAS_ID=<your-bedrock-agent-alias-id> OPENAI_API_KEY=<your-openai-api-key>
OPENAI_ASSISTANT_ID=<your-openai-assistant-id>
```

- Create Dockerfile:
  - `vi Dockerfile`
  - Content of Dockerfile:

```
FROM node:18
WORKDIR /usr/src/app
COPY package*.json ./
RUN npm install
COPY . .
EXPOSE 5000
CMD ["npm", "start"]
```

### Frontend

- Create folder and download source code:

```
cd ..
mkdir frontend && cd frontend
```

wget <https://tcb-public-events.s3.amazonaws.com/mdac/resources/day2/cloudmart-frontend.zip> unzip cloudmart-frontend.zip

- Create Dockerfile:
  - vi Dockerfile
  - Content of Dockerfile:

```
FROM node:16-alpine as build
WORKDIR /app
COPY package*.json ./
RUN npm ci
COPY . .
RUN npm run build

FROM node:16-alpine
WORKDIR /app
RUN npm install -g serve
COPY --from=build /app/dist /app
ENV PORT=5001
ENV NODE_ENV=production
EXPOSE 5001
CMD ["serve", "-s", ".", "-l", "5001"]
```

## Part 2 – Kubernetes

Attention: AWS Kubernetes service is not free, so when executing the hands-on below, you will be charged a few cents on your AWS account according to EKS pricing on AWS.

Remember to delete the cluster to avoid unwanted charges. Use the removal section at the end of the doc.

### Cluster Setup on AWS Elastic Kubernetes Services (EKS)

1. Create a user named eksuser with Admin privileges and authenticate with it

```
aws configure
```

2. Install the CLI tool eksctl

```
curl --silent --location https://github.com/weaveworks/eksctl/releases/latest/download/eksctl $(uname -s) _amd64.tar.gz | tar xz -C /tmp
sudo cp /tmp/eksctl /usr/bin
eksctl version
```

### 3. Install the CLI tool kubectl

```
curl -o kubectl https://amazon-eks.s3.us-west-2.amazonaws.com/1.18.9/2020-11-02/bin/linux/amd64/kubectl
chmod +x ./kubectl
mkdir -p $HOME/bin && cp ./kubectl $HOME/bin/kubectl && export PATH=$PATH:$HOME/bin
echo 'export PATH=$PATH:$HOME/bin' >> ~/.bashrc
kubectl version --short --client
```

### 4. Create an EKS Cluster

```
eksctl create cluster \
  --name cloudmart \
  --region us-east-1 \
  --nodegroup-name standard-workers \
  --node-type t3.medium \
  --nodes 1 \
  --with-oidc \
  --managed
```

### 5. Connect to the EKS cluster using the kubectl configuration

```
aws eks update-kubeconfig --name cloudmart
```

### 6. Verify Cluster Connectivity

```
kubectl get svc
kubectl get nodes
```

### 7. Create a Role & Service Account to provide pods access to services used by the application (DynamoDB, Bedrock, etc).

```
eksctl create iamserviceaccount \
  --cluster=cloudmart \
  --name=cloudmart-pod-execution-role \
  --role-name CloudMartPodExecutionRole \
  --attach-policy-arn=arn:aws:iam::aws:policy/AdministratorAccess \
  --region us-east-1 \
  --approve
```

**NOTE:** In the example above, Admin privileges were used to facilitate educational purposes. Always remember to follow the principle of least privilege in production environments

## Backend Deployment on Kubernetes

- Create an ECR Repository for the Backend and upload the Docker image to it  
Repository name: cloudmart-backend (do this step in AWS)
- Switch to backend folder

```
cd ../../
cd challenge-day2/backend
```

- Follow the ECR steps to build your Docker image( As shown in the ECR Repository)
- Create a Kubernetes deployment file (YAML) for the Backend

```
vi cloudmart-backend.yaml
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: cloudmart-backend-app
spec:
  replicas: 1
  selector:
    matchLabels:
      app: cloudmart-backend-app
  template:
    metadata:
      labels:
        app: cloudmart-backend-app
    spec:
      serviceAccountName: cloudmart-pod-execution-role
      containers:
        - name: cloudmart-backend-app
          image: public.ecr.aws/l4c0j8h9/cloudmart-backend:latest
          env:
            - name: PORT
              value: "5000"
            - name: AWS_REGION
              value: "us-east-1"
            - name: BEDROCK_AGENT_ID
              value: "xxxxxx"
            - name: BEDROCK_AGENT_ALIAS_ID
              value: "xxxx"
            - name: OPENAI_API_KEY
              value: "xxxxxx"
            - name: OPENAI_ASSISTANT_ID
              value: "xxxx"
---
apiVersion: v1
kind: Service
metadata:
  name: cloudmart-backend-app-service
spec:
  type: LoadBalancer
  selector:
    app: cloudmart-backend-app
  ports:
    - protocol: TCP
      port: 5000
      targetPort: 5000
```

- Deploy the Backend on Kubernetes

```
kubectl apply -f cloudmart-backend.yaml
```

- Monitor the status of objects being created and obtain the public IP generated for the API

```
kubectl get pods
kubectl get deployment
kubectl get service
```

# Frontend Deployment on Kubernetes

## Preparation

- Change the Frontend's .env file to point to the API URL created within Kubernetes obtained by the kubectl get service command

```
cd ../challenge-day2/frontend  
nano .env
```

Content of .env:

```
VITE_API_BASE_URL=http://<your_url_kubernetes_api>:5000/api
```

- Create an ECR Repository for the Frontend and upload the Docker image to it
  - Repository name: cloudmart-frontend (Do this step in AWS).
- Follow the ECR steps to build your Docker image ( As shown in the AWS ECR).
- Create a Kubernetes deployment file (YAML) for the Frontend

```
nano cloudmart-frontend.yaml
```

```
apiVersion: apps/v1  
kind: Deployment  
metadata:  
  name: cloudmart-frontend-app  
spec:  
  replicas: 1  
  selector:  
    matchLabels:  
      app: cloudmart-frontend-app  
  template:  
    metadata:  
      labels:  
        app: cloudmart-frontend-app  
    spec:  
      serviceAccountName: cloudmart-pod-execution-role  
      containers:  
        - name: cloudmart-frontend-app  
          image: public.ecr.aws/l4c0j8h9/cloudmart-frontend:latest  
---  
apiVersion: v1  
kind: Service  
metadata:  
  name: cloudmart-frontend-app-service  
spec:  
  type: LoadBalancer  
  selector:  
    app: cloudmart-frontend-app  
  ports:  
    - protocol: TCP  
      port: 5001  
      targetPort: 5001
```

- Deploy the Frontend on Kubernetes

```
kubectl apply -f cloudmart-frontend.yaml
```

- Monitor the status of objects being created and obtain the public IP generated for the API

```
kubectl get pods  
kubectl get deployment  
kubectl get service
```

## Removal

- At the end of the hands-on, delete all resources:
- If you delete the cluster at the end of the exercise, you'll have to recreate it for the next days. So decide what makes more sense for you: delete the cluster and recreate it every day or keep it and pay for the time it's running. However, don't forget to delete it permanently at the end of the Challenge.

```
kubectl delete service cloudmart-frontend-app-service
kubectl delete deployment cloudmart-frontend-app
kubectl delete service cloudmart-backend-app-service
kubectl delete deployment cloudmart-backend-app

eksctl delete cluster --name cloudmart --region us-east-1
```